

## Designing a Scheduler.

### Task Setup:

ID	Description	Duration	Dependencies	Status	Type
0	Get ready for the day (wake up, get dressed, make breakfast, go to cafe)	60	None	not_yet_started	Flexible
1	PCW	120	0	not_yet_started	Flexible
2	NS144 Class	90	1	not_yet_started	Fixed
3	Lunch	60	2	not_yet_started	Flexible
4	Work	60	3	not_yet_started	Flexible
5	Commute to Lab	60	4	not_yet_started	Flexible
6	Lab Meeting	120	5	not_yet_started	Fixed
7	Groceries (+ commute back)	60	6	not_yet_started	Flexible
8	Dinner (Steak because we are in Argentina)	60	7	not_yet_started	Flexible
9	Get ready for sleep	60	8	not_yet_started	Flexible

Total Scheduler Time: 15 h (900 minutes)

### [High Level Algorithmic Strategy Video](#)

### Python Implementation:

```

Scheduler:
Start:07:00 -> Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 08:00 -> Completed 'Get ready (wake up, dress, breakfast, café)'

-----
Start:08:00 -> Starting 'PCW (morning)' (priority=38.0)
End: 10:00 -> Completed 'PCW (morning)'

-----
Start:10:00 -> Starting 'NS144 Class' (priority=76.0)
End: 11:30 -> Completed 'NS144 Class'

-----
Start:11:30 -> Starting 'Lunch' (priority=36.0)
End: 13:00 -> Completed 'Lunch'

-----
Start:13:00 -> Starting 'Work' (priority=39.0)
End: 14:00 -> Completed 'Work'

-----
Start:14:00 -> Starting 'Commute to Lab' (priority=39.0)
End: 15:00 -> Completed 'Commute to Lab'

-----
Start:15:00 -> Starting 'Lab time' (priority=73.0)
End: 17:00 -> Completed 'Lab time'

-----
Start:17:00 -> Starting 'Groceries (+ commute back)' (priority=39.0)
End: 18:00 -> Completed 'Groceries (+ commute back)'

-----
Start:18:00 -> Starting 'Dinner (steak)' (priority=39.0)
End: 19:00 -> Completed 'Dinner (steak)'

-----
Start:19:00 -> Starting 'PCW CS 110' (priority=33.0)
End: 21:00 -> Completed 'PCW CS 110'

-----
Start:21:00 -> Starting 'Get ready for sleep' (priority=39.0)
End: 22:00 -> Completed 'Get ready for sleep'

-----
All tasks are done in 900 minutes.

```

*Figure 1: Output of Scheduler Algorithm.*

**Is the schedule that your code returns feasible? (E.g., having breakfast at 11pm might not be ideal). Do you anticipate running into trouble while test-driving this code? Briefly explain in no more than 100 words.**

Yes, I think the schedule my code returns is mostly feasible. The output shows tasks in a logical order and correctly computes total runtime. This makes sense since fixed tasks like NS144 class (priority 76) and Lab (73) are given higher priority, while each task depends on the one before it, which naturally enforces a realistic sequence. However, the schedule does not yet account for real-world timing, such as breaks, commuting, or unforeseen circumstances. This means small delays could “cascade” and push later tasks into other time frames.

**Algorithm Test-Drive:**  [CS 110 Scheduler Test-Drive.mov](#)

[Algorithmic Analyses Video](#)

**The plots generated:**

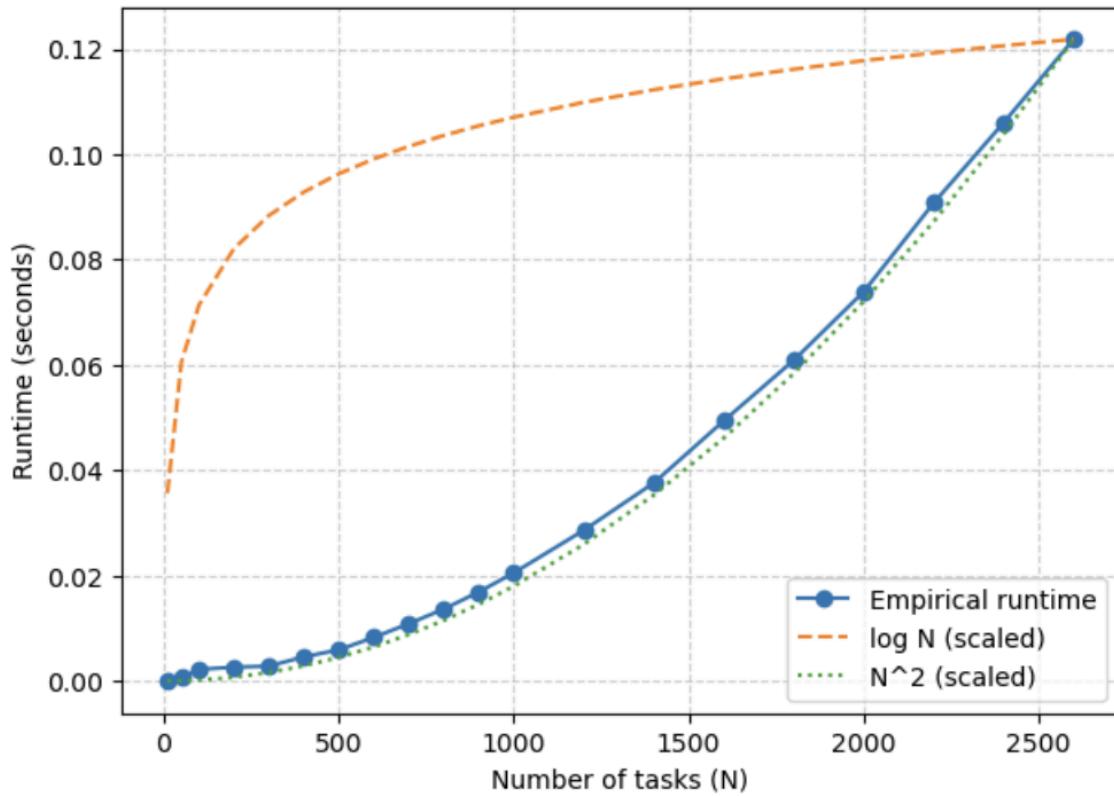
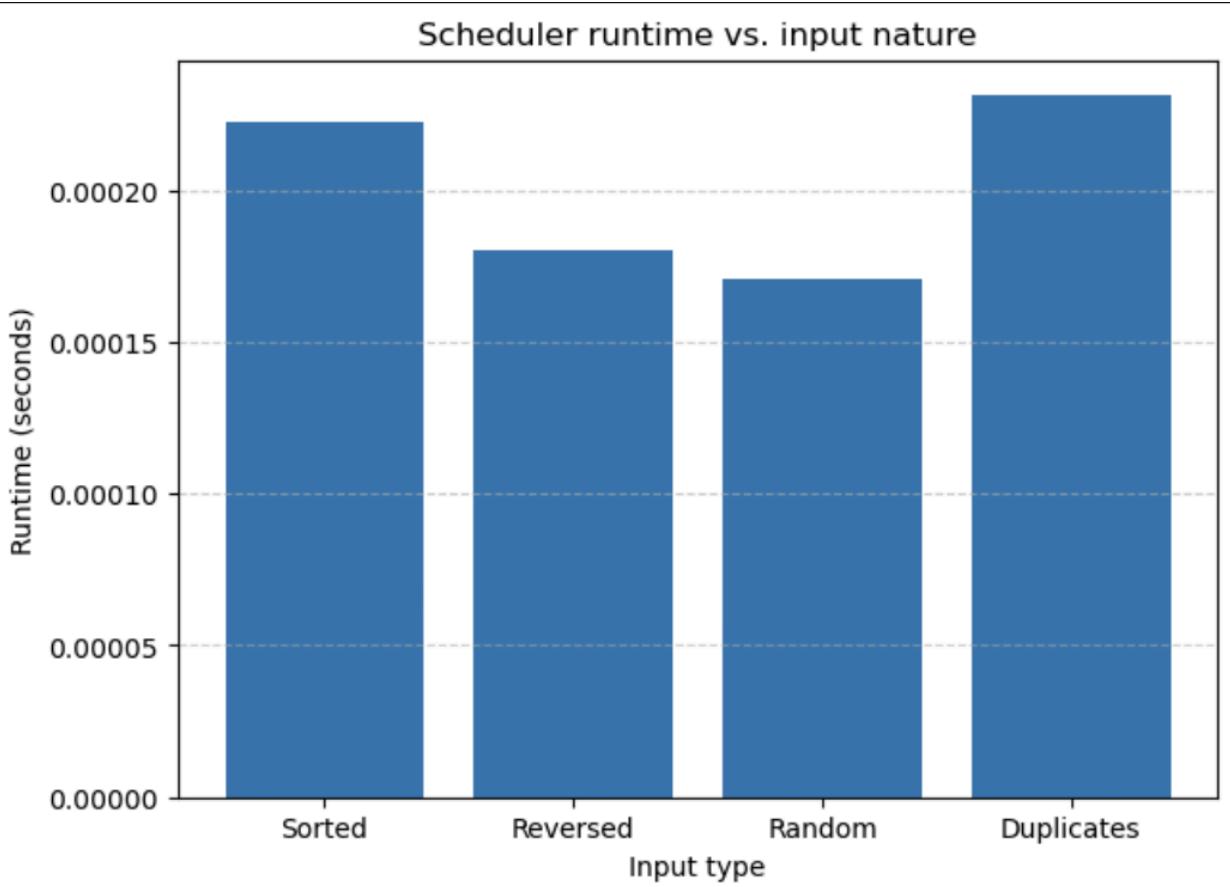


Figure 2: Input size scaling.



*Figure 3: Input nature (order) runtime.*

**Any specific calculations that supported your argument:**

Scheduler.run() loop:

Even though heap operations are  $O(\log N)$ , the dominant cost actually comes from the body of the main loop:

```
# Roughly one iteration per completed task

for task in self.tasks:                      # O(N)

    if task.status == 'not_yet_started':        # O(1)

        and self.all_deps_completed(task):      # small, depends on #deps

        and task not in self.heap.heap:         # O(N) membership check

        self.heap.heappush(task)                # O(log N)
```

```
task.status = 'in_progress'
```

Full scan of all tasks =  $O(N)$ : Each iteration of the loop checks every task to see if it's ready to run. This means that for every task completed, we still iterate through the entire list of tasks.

Linear membership check =  $O(N)$ : The condition task not in `self.heap.heap` checks membership in a Python list, which takes  $O(N)$  time. This adds another linear cost per iteration.

Heap operations =  $O(\log N)$ : Inserting a new task (`heappush`) is logarithmic, but its contribution is much smaller than the repeated linear scans.

Since the main while loop runs approximately  $N$  times (once per completed task), these operations multiply together:

$$T(N) = N * [O(N) + O(N) + O(\log N)] \sim O(N^2) \rightarrow \text{In this case the dominant term is } O(n^2)$$

#### **A written summary of results (<100 words):**

The runtime of the scheduler increases  $\sim$  quadratically as the number of tasks grows. While heap operations are  $O(\log N)$ , the repeated scanning of all tasks and linear membership checks dominate the total cost, resulting in  $O(N^2)$  scaling. The empirical curve in the plot closely matches the  $N^2$  line, confirming that as input size increases, runtime grows much faster than logarithmic scaling would predict. This means the scheduler performs efficiently for small task sets but becomes increasingly slow for large ones due to repeated full-list operations within each scheduling loop.

#### **Improvement:**

**Speed:** I can combat the  $O(N^2)$  behavior by keeping dependencies in a dict of sets ( $O(1)$  add/remove/lookup), tracking a live “ready” set instead of rescanning everything each loop, and swapping task not in heap list checks for a companion set so membership is  $O(1)$ .

**Feasibility:** Add small buffers between tasks, soft time windows (e.g., “eat before 9pm”), and handle cascading delays. I’d also make durations commute-aware and reserve a bit of slack for long blocks so one overrun doesn’t derail the whole evening..

Word count: 89

#### **Reflection Video**

## Appendix:

[Canva Whiteboard](#)

### Python Code

ID	Description	Duration (min)	Dependencies	Status	Type
0	Get ready for the day (wake up, get dressed, make breakfast, go to café)	60	None	not_yet_started	Flexible
1	PCW	120	0	not_yet_started	Flexible
2	NS144 Class	90	1	not_yet_started	Fixed
3	Lunch	90	2	not_yet_started	Flexible
4	Work	60	3	not_yet_started	Flexible
5	Commute to Lab	60	4	not_yet_started	Flexible
6	Lab Meeting	120	5	not_yet_started	Fixed
7	Groceries (+ commute back)	60	6	not_yet_started	Flexible
8	Dinner (Steak because we are in Argentina)	60	7	not_yet_started	Flexible
9	PCW CS 110	120	8	not_yet_started	Flexible
10	Get ready for sleep	60	9	not_yet_started	Flexible

```
[110]: # 1. Task class

class Task:
    """
    Represents a single task in the scheduler.

    Attributes:
        id (int): Unique identifier for the task
        description (str): Short description of the task
        duration (int): Duration of the task in minutes
        dependencies (list[int]): List of task IDs that must be completed first, before task can be completed
        status (str): Current status ('not_yet_started', 'in_progress', 'completed')
        type (str): 'Fixed' or 'Flexible'
        priority (float): Automatically computed priority value

    creating this class will be useful to build on the input data later - it lets me easily compute task properties,
    manage dependencies, and update attributes like status or priority as the scheduler runs.
    """

    def __init__(self, id, description, duration, dependencies=None, status='not_yet_started', type='Flexible'):
        self.id = id
        self.description = description
        self.duration = duration
        # handles both cases (tasks with/without dependencies)
        if dependencies:
            self.dependencies = [int(d) for d in str(dependencies).replace(" ", "").split(",")]
        else:
            self.dependencies = []

        self.status = status
        self.type = type
        self.priority = 0.0 # will be computed later based on task type, duration, and dependencies

    def __repr__(self):
        return Task(self.id), {self.description}, priority=(self.priority:.1))
```

```

# 2. Create MaxHeap class to compute priority queue

class MaxHeap:
    """
    Implementation of a MaxHeap (priority queue).
    The highest-priority element is always at index 0.

    Methods:
        left(i), right(i), parent(i): return child or parent index positions
        heappush(task): insert a new task and restore heap order
        heappop(): remove and return the highest-priority task
        heapify(i): restore heap property starting from index i

    This class will be used in the scheduler algorithm to select the next task
    with the highest priority. The heap helps maintain the heap property
    (parents>children) after every insertion or removal, giving the O(log n)
    performance for heappush and heappop operations.
    """

    def __init__(self):
        self.heap = []

    def left(self, i): return 2 * i + 1
    def right(self, i): return 2 * i + 2
    def parent(self, i): return (i - 1) // 2

    def heappush(self, task):
        """add a new task and move it up until the heap property is restored."""
        self.heap.append(task)
        index = len(self.heap) - 1
        # while parent has smaller priority, swap upward
        while index > 0:
            parent = self.parent(index)
            if self.heap[parent].priority >= self.heap[index].priority:
                break
            self.heap[parent], self.heap[index] = self.heap[index], self.heap[parent]
            index = parent

    def heapify(self, i):
        """Keep the largest priority task on top (restores heap order)."""
        size = len(self.heap)
        left = self.left(i)
        right = self.right(i)
        largest = i

        # check left child
        if left < size and self.heap[left].priority > self.heap[largest].priority:
            largest = left

        # check right child
        if right < size and self.heap[right].priority > self.heap[largest].priority:
            largest = right

        # swap if parent is smaller than a child
        if largest != i:
            self.heap[i], self.heap[largest] = self.heap[largest], self.heap[i]
            self.heapify(largest) # keep fixing down the tree

    def heappop(self):
        """Remove and return the task with the highest priority"""
        n = len(self.heap)
        if n == 0:
            return None #base case
        if n == 1:
            return self.heap.pop()

        max_item = self.heap[0] # save root (max)
        self.heap[0] = self.heap.pop() # move last to root
        self.heapify(0) # restore heap property
        return max_item

    def __len__(self):
        return len(self.heap)

```

```
# 3. Scheduler class

class Scheduler:
    """
    Schedules and executes a list of tasks based on priority and dependencies.
    1) It takes a list of Task objects with attributes such as duration, dependencies, and type
    2) It uses MaxHeap to pick the next highest-priority task
    3) Makes sure dependencies are completed before execution
    4) Keeps track of time (in minutes)

    Attributes:
        tasks : list[Task] - All Task objects to be scheduled.
        current_time : int - clock in minutes, default 7*60 = 7:00 AM
        completed : list[int] - IDs of tasks that have been finished, for "unlocking" dependent tasks
        heap : MaxHeap - priority queue to store all ready tasks, keeping the highest priority at the top.
    """

    def __init__(self, tasks, start_time=7*60):
        self.tasks = tasks
        self.current_time = start_time # in minutes (7*60 = 7:00 AM)
        self.completed = []
        self.heap = MaxHeap()

    def compute_priority(self, task):
        """
        Compute a priority value: higher for fixed tasks, fewer dependencies, and shorter duration
        Logic:
        - 50 base score
        - +40 if task.type == Fixed (time-sensitive fixed task)
        - -5 for each dependency (because it makes it harder to start the task)
        - -duration/10 because we prefer shorter tasks that keeps us move forward
        """
        base_score = 50 #starting point

        if str(task.type).lower() == "fixed":
            base_score += 40 #time sensitive boost for fixed tasks (class, meeting)

        dependency_penalty = 5*len(task.dependencies)
        duration_penalty = task.duration / 10

        #computing final priority value
        priority = base_score - dependency_penalty - duration_penalty
        return priority

    def all_deps_completed(self, task):
        """
        Check if all dependencies of a task are completed.
        """
        return all(dep in self.completed for dep in task.dependencies)

    def run(self):
        """
        Run the scheduler and print the execution plan + time
        Logic (priority queue operations)
        1) Compute priority for each task (like assigning a key value in a heap)
        2) Check tasks that are ready = dependencies are complete
        3) Add ready tasks in MaxHeap
        4) Pop task with highest priority and do it
        5) Mark task as done and go onto tasks that where dependent on it
        6) Algorithms stops when there are no more tasks to perform

        Time complexity: heappush() and heappop() runs in O(log n) as described in Cormen et al.
        """
        # Compute the priority for each task
        for task in self.tasks:
            task.priority = self.compute_priority(task)

        print("Scheduler:")# start the scheduler

        while len(self.completed) < len(self.tasks): #keeps looping until all tasks are completed
            # add ready (with no dependencies) tasks to heap
            for task in self.tasks:
                if (task.status == 'not_yet_started' and
                    self.all_deps_completed(task) and
                    task not in self.heap.heap):
                    self.heap.heappush(task)
                task.status = 'in_progress' #update status

            # gets the next tasks with highest priority (root node)
            next_task = self.heap.heappop()
            if not next_task:
                print("no task ready")
                break

            # Execute the task
            start_hr = self.current_time // 60 #time in hours (for time stamps)
            start_min = self.current_time % 60 #time in minutes for the total duration of scheduler

            # using f-strings to allow including the variables
            print(f"Start:{start_hr:02d}:{start_min:02d} - Starting '{next_task.description}' (priority={next_task.priority:.1f})")

            self.current_time += next_task.duration #add time on the clock
            next_task.status = 'completed' #update status
            self.completed.append(next_task.id) #go on to the next task

            end_hr = self.current_time // 60
            end_min = self.current_time % 60
            print(f"End: {end_hr:02d}:{end_min:02d} - Completed '{next_task.description}'")

            #adding a break between each completed task: easier to read output
            print("-----")

        print(f"All tasks are done in {self.current_time - 7*60} minutes.")
```

```
[112]: # Example 1: ordered list of tasks

tasks = [
    Task(0, "Get ready (wake up, dress, breakfast, café)", 60, None, 'not_yet_started', 'Flexible'),
    Task(1, "PCW (morning)", 120, 0, 'not_yet_started', 'Flexible'),
    Task(2, "NS144 Class", 90, 1, 'not_yet_started', 'Fixed'),
    Task(3, "Lunch", 90, 2, 'not_yet_started', 'Flexible'),
    Task(4, "Work", 60, 3, 'not_yet_started', 'Flexible'),
    Task(5, "Commute to Lab", 60, 4, 'not_yet_started', 'Flexible'),
    Task(6, "Lab time", 120, 5, 'not_yet_started', 'Fixed'),
    Task(7, "Groceries (+ commute back)", 60, 6, 'not_yet_started', 'Flexible'),
    Task(8, "Dinner (steak)", 60, 7, 'not_yet_started', 'Flexible'),
    Task(9, "PCW CS 110", 120, 8, 'not_yet_started', 'Flexible'),
    Task(10, "Get ready for sleep", 60, 9, 'not_yet_started', 'Flexible'),
]

# Run the scheduler
my_scheduler_ordered = Scheduler(tasks)
my_scheduler_ordered.run()
tasks_sorted = tasks

Scheduler:
Start:07:00 -> Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 08:00 -> Completed 'Get ready (wake up, dress, breakfast, café)'

Start:08:00 -> Starting 'PCW (morning)' (priority=38.0)
End: 10:00 -> Completed 'PCW (morning)'

Start:10:00 -> Starting 'NS144 Class' (priority=76.0)
End: 11:30 -> Completed 'NS144 Class'

Start:11:30 -> Starting 'Lunch' (priority=36.0)
End: 13:00 -> Completed 'Lunch'

Start:13:00 -> Starting 'Work' (priority=39.0)
End: 14:00 -> Completed 'Work'

Start:14:00 -> Starting 'Commute to Lab' (priority=39.0)
End: 15:00 -> Completed 'Commute to Lab'

Start:15:00 -> Starting 'Lab time' (priority=73.0)
End: 17:00 -> Completed 'Lab time'

Start:17:00 -> Starting 'Groceries (+ commute back)' (priority=39.0)
End: 18:00 -> Completed 'Groceries (+ commute back)'

Start:18:00 -> Starting 'Dinner (steak)' (priority=39.0)
End: 19:00 -> Completed 'Dinner (steak)'

Start:19:00 -> Starting 'PCW CS 110' (priority=33.0)
End: 21:00 -> Completed 'PCW CS 110'

Start:21:00 -> Starting 'Get ready for sleep' (priority=39.0)
End: 22:00 -> Completed 'Get ready for sleep'

All tasks are done in 900 minutes.
```

```
[116]: # Example 2: reversed ordered list of tasks (worst case)

tasks = [
    Task(10, "Get ready for sleep", 60, 9, 'not_yet_started', 'Flexible'),
    Task(9, "PCW CS 110", 120, 8, 'not_yet_started', 'Flexible'),
    Task(8, "Dinner (steak)", 60, 7, 'not_yet_started', 'Flexible'),
    Task(7, "Groceries (+ commute back)", 60, 6, 'not_yet_started', 'Flexible'),
    Task(6, "Lab time", 120, 5, 'not_yet_started', 'Fixed'),
    Task(5, "Commute to Lab", 60, 4, 'not_yet_started', 'Flexible'),
    Task(4, "Work", 60, 3, 'not_yet_started', 'Flexible'),
    Task(3, "Lunch", 90, 2, 'not_yet_started', 'Flexible'),
    Task(2, "NS144 Class", 90, 1, 'not_yet_started', 'Fixed'),
    Task(1, "PCW (morning)", 120, 0, 'not_yet_started', 'Flexible'),
    Task(0, "Get ready (wake up, dress, breakfast, café)", 60, None, 'not_yet_started', 'Flexible'),
]

# Run the scheduler
my_scheduler_reversed = Scheduler(tasks)
my_scheduler_reversed.run()
tasks_reversed = tasks

Scheduler:
Start:07:00 → Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 08:00 → Completed 'Get ready (wake up, dress, breakfast, café)'

-----
Start:08:00 → Starting 'PCW (morning)' (priority=38.0)
End: 10:00 → Completed 'PCW (morning)'

-----
Start:10:00 → Starting 'NS144 Class' (priority=76.0)
End: 11:30 → Completed 'NS144 Class'

-----
Start:11:30 → Starting 'Lunch' (priority=36.0)
End: 13:00 → Completed 'Lunch'

-----
Start:13:00 → Starting 'Work' (priority=39.0)
End: 14:00 → Completed 'Work'

-----
Start:14:00 → Starting 'Commute to Lab' (priority=39.0)
End: 15:00 → Completed 'Commute to Lab'

-----
Start:15:00 → Starting 'Lab time' (priority=73.0)
End: 17:00 → Completed 'Lab time'

-----
Start:17:00 → Starting 'Groceries (+ commute back)' (priority=39.0)
End: 18:00 → Completed 'Groceries (+ commute back)'

-----
Start:18:00 → Starting 'Dinner (steak)' (priority=39.0)
End: 19:00 → Completed 'Dinner (steak)'

-----
Start:19:00 → Starting 'PCW CS 110' (priority=33.0)
End: 21:00 → Completed 'PCW CS 110'

-----
Start:21:00 → Starting 'Get ready for sleep' (priority=39.0)
End: 22:00 → Completed 'Get ready for sleep'

-----
All tasks are done in 900 minutes.
```

```
[118]: # Example 3: randomly ordered list of tasks (test for robustness)

tasks = [
    Task(4, "Work", 60, 3, 'not_yet_started', 'Flexible'),
    Task(9, "PCW CS 110", 120, 8, 'not_yet_started', 'Flexible'),
    Task(0, "Get ready (wake up, dress, breakfast, café)", 60, None, 'not_yet_started', 'Flexible'),
    Task(6, "Lab time", 120, 5, 'not_yet_started', 'Fixed'),
    Task(2, "NS144 Class", 90, 1, 'not_yet_started', 'Fixed'),
    Task(10, "Get ready for sleep", 60, 9, 'not_yet_started', 'Flexible'),
    Task(5, "Commute to Lab", 60, 4, 'not_yet_started', 'Flexible'),
    Task(3, "Lunch", 90, 2, 'not_yet_started', 'Flexible'),
    Task(1, "PCW (morning)", 120, 0, 'not_yet_started', 'Flexible'),
    Task(8, "Dinner (steak)", 60, 7, 'not_yet_started', 'Flexible'),
    Task(7, "Groceries (+ commute back)", 60, 6, 'not_yet_started', 'Flexible'),
]

# Run the scheduler
my_scheduler_random = Scheduler(tasks)
my_scheduler_random.run()
tasks_random = tasks

Scheduler:
Start:07:00 → Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 08:00 → Completed 'Get ready (wake up, dress, breakfast, café)'
-----
Start:08:00 → Starting 'PCW (morning)' (priority=38.0)
End: 10:00 → Completed 'PCW (morning)'
-----
Start:10:00 → Starting 'NS144 Class' (priority=76.0)
End: 11:30 → Completed 'NS144 Class'
-----
Start:11:30 → Starting 'Lunch' (priority=36.0)
End: 13:00 → Completed 'Lunch'
-----
Start:13:00 → Starting 'Work' (priority=39.0)
End: 14:00 → Completed 'Work'
-----
Start:14:00 → Starting 'Commute to Lab' (priority=39.0)
End: 15:00 → Completed 'Commute to Lab'
-----
Start:15:00 → Starting 'Lab time' (priority=73.0)
End: 17:00 → Completed 'Lab time'
-----
Start:17:00 → Starting 'Groceries (+ commute back)' (priority=39.0)
End: 18:00 → Completed 'Groceries (+ commute back)'
-----
Start:18:00 → Starting 'Dinner (steak)' (priority=39.0)
End: 19:00 → Completed 'Dinner (steak)'
-----
Start:19:00 → Starting 'PCW CS 110' (priority=33.0)
End: 21:00 → Completed 'PCW CS 110'
-----
Start:21:00 → Starting 'Get ready for sleep' (priority=39.0)
End: 22:00 → Completed 'Get ready for sleep'
-----
All tasks are done in 900 minutes.
```

```
[120]: # Example 4: list with duplicate tasks (edge case test)

tasks = [
    Task(0, "Get ready (wake up, dress, breakfast, café)", 60, None, 'not_yet_started', 'Flexible'),
    Task(1, "PCW (morning)", 120, 0, 'not_yet_started', 'Flexible'),
    Task(2, "NS144 Class", 90, 1, 'not_yet_started', 'Fixed'),
    Task(3, "Lunch", 90, 2, 'not_yet_started', 'Flexible'),
    Task(4, "Work", 60, 3, 'not_yet_started', 'Flexible'),
    Task(5, "Commute to Lab", 60, 4, 'not_yet_started', 'Flexible'),
    Task(6, "Lab time", 120, 5, 'not_yet_started', 'Fixed'),
    Task(7, "Groceries (+ commute back)", 60, 6, 'not_yet_started', 'Flexible'),
    Task(8, "Dinner (steak)", 60, 7, 'not_yet_started', 'Flexible'),
    Task(9, "PCW CS 110", 120, 8, 'not_yet_started', 'Flexible'),
    Task(10, "Get ready for sleep", 60, 9, 'not_yet_started', 'Flexible'),

    # duplicates, same id/description
    Task(3, "Lunch", 90, 2, 'not_yet_started', 'Flexible'), # same ID and description
    Task(6, "Lab time", 120, 5, 'not_yet_started', 'Fixed'), # same ID again
    Task(11, "Get ready (wake up, dress, breakfast, café)", 60, None, 'not_yet_started', 'Flexible'), # same description, new ID
]

# Run the scheduler
my_scheduler_duplicates = Scheduler(tasks)
my_scheduler_duplicates.run()
tasks_duplicates = tasks

Scheduler:
Start:07:00 -> Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 08:00 -> Completed 'Get ready (wake up, dress, breakfast, café)'
-----
Start:08:00 -> Starting 'Get ready (wake up, dress, breakfast, café)' (priority=44.0)
End: 09:00 -> Completed 'Get ready (wake up, dress, breakfast, café)'
-----
Start:09:00 -> Starting 'PCW (morning)' (priority=38.0)
End: 11:00 -> Completed 'PCW (morning)'
-----
Start:11:00 -> Starting 'NS144 Class' (priority=76.0)
End: 12:30 -> Completed 'NS144 Class'
-----
Start:12:30 -> Starting 'Lunch' (priority=36.0)
End: 14:00 -> Completed 'Lunch'
-----
Start:14:00 -> Starting 'Work' (priority=39.0)
End: 15:00 -> Completed 'Work'
-----
Start:15:00 -> Starting 'Commute to Lab' (priority=39.0)
End: 16:00 -> Completed 'Commute to Lab'
-----
Start:16:00 -> Starting 'Lab time' (priority=73.0)
End: 18:00 -> Completed 'Lab time'
-----
Start:18:00 -> Starting 'Lab time' (priority=73.0)
End: 20:00 -> Completed 'Lab time'
-----
Start:20:00 -> Starting 'Groceries (+ commute back)' (priority=39.0)
End: 21:00 -> Completed 'Groceries (+ commute back)'
-----
Start:21:00 -> Starting 'Dinner (steak)' (priority=39.0)
End: 22:00 -> Completed 'Dinner (steak)'
-----
Start:22:00 -> Starting 'Lunch' (priority=36.0)
End: 23:30 -> Completed 'Lunch'
-----
Start:23:30 -> Starting 'PCW CS 110' (priority=33.0)
End: 25:30 -> Completed 'PCW CS 110'
-----
Start:25:30 -> Starting 'Get ready for sleep' (priority=39.0)
End: 26:30 -> Completed 'Get ready for sleep'

All tasks are done in 1170 minutes.
```

```

Part 2: Efficiency of the scheduler

[122]: # Empirical Efficiency
# ( How many tasks I did (during the scheduled time frame) / How many tasks were scheduled ) * 100

print("Percentage of tasks achieved (during the scheduled time frame):", (7/11) * 100)

Percentage of tasks achieved (during the scheduled time frame): 63.63636363636363

Experimental Design:
1) Input nature (order) scaling - How runtime changes based on input nature?
2) Input size scaling - How runtime changes as we increase task count?

[130]: #1) Input nature (order) scaling - How runtime changes based on input nature?
# measure how long the scheduler takes for each input order

# minimal timing helper
import time, io, contextlib

def time_scheduler(task_list):
    """Reset tasks, run Scheduler once, return seconds (prints muted)."""
    # reset simple state
    for t in task_list:
        t.status = 'not_yet_started'
        t.priority = 0.0
    s = Scheduler(task_list)
    # silence the scheduler prints so timing isn't noisy
    buf = io.StringIO()
    with contextlib.redirect_stdout(buf):
        start = time.perf_counter()
        s.run()
        end = time.perf_counter()
    return end - start

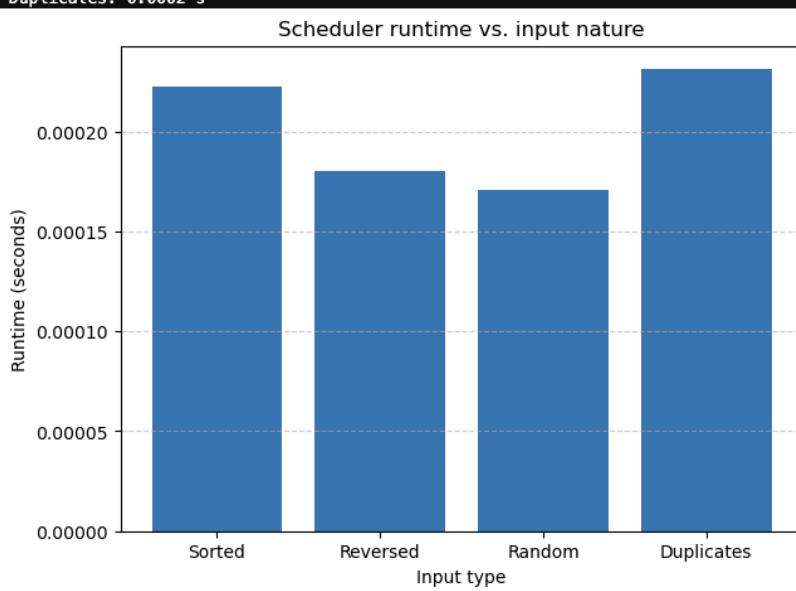
labels = ["Sorted", "Reversed", "Random", "Duplicates"]
inputs = [tasks_sorted, tasks_reversed, tasks_random, tasks_duplicates]

runtimes = []
for name, lst in zip(labels, inputs):
    secs = time_scheduler(lst)
    runtimes.append(secs)
    print(f"{name}: {secs:.4f} s")

plt.figure(figsize=(7,5))
plt.bar(labels, runtimes)
plt.title("Scheduler runtime vs. input nature")
plt.xlabel("Input type")
plt.ylabel("Runtime (seconds)")
plt.grid(axis="y", linestyle="--", alpha=0.6)
plt.show()

```

Sorted: 0.0002 s  
 Reversed: 0.0002 s  
 Random: 0.0002 s  
 Duplicates: 0.0002 s



```
[315]: # 2) Input size scaling - How runtime changes as we increase task count?
# Plan:
#   - Build N tasks (no dependencies)
#   - Time Scheduler.run() for each N
#   - simulate a few times and average
#   - Plot N vs runtime

import time, io, contextlib
import matplotlib.pyplot as plt
import random

#create N tasks (no dependencies because input nature doesn't matter here)
def make_tasks(n):
    tasks = []
    for i in range(n):
        duration = random.randint(10, 60)      # small variation
        ttype = 'Fixed' if (i % 5 == 0) else 'Flexible'
        tasks.append(Task(i, f"Task {i}", duration, None, 'not_yet_started', ttype))
    return tasks

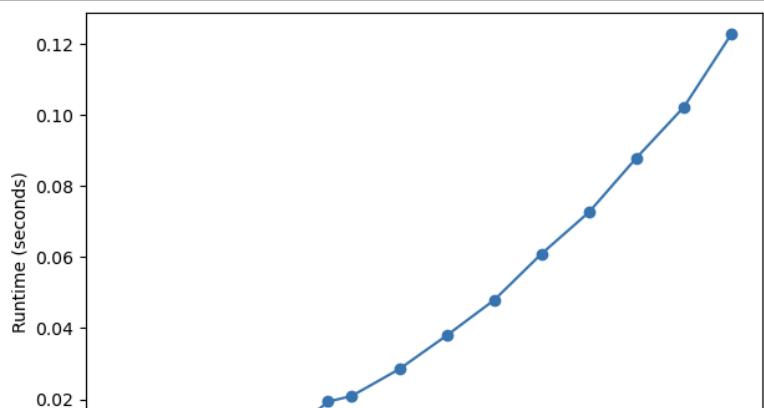
# Measure average runtime over a few repeats (reduces randomness)
def avg_runtime(task_list, repeats=5):
    total = 0.0
    for _ in range(repeats):
        total += time_scheduler_quiet(task_list)
    return total / repeats

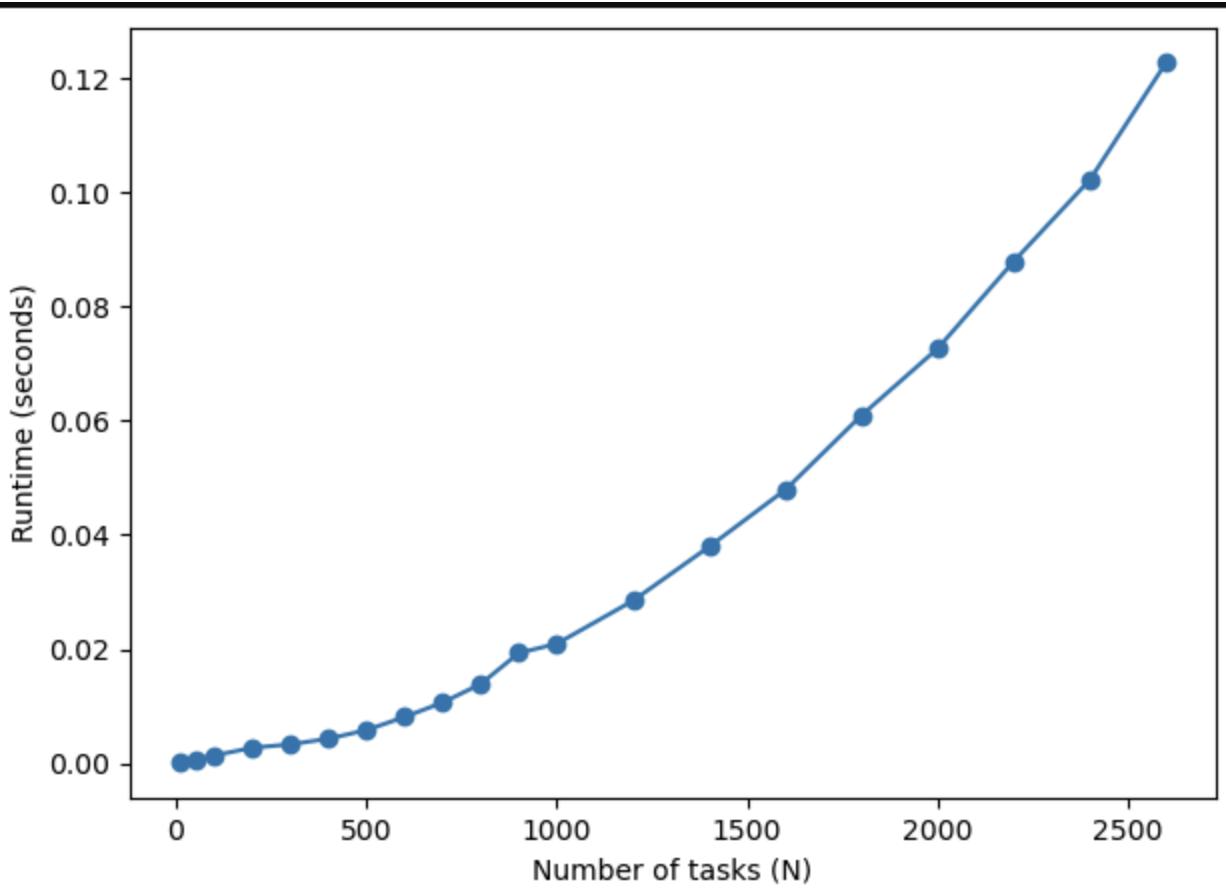
sizes = [10, 50, 100, 200, 300, 400, 500, 600, 700, 800, 900, 1000, 1200, 1400, 1600, 1800, 2000, 2200, 2400, 2600]
runtimes = []

print("Input size scaling (avg of 10 runs each task):")
for n in sizes:
    tasksN = make_tasks(n)
    rt = avg_runtime(tasksN, repeats=10)
    runtimes.append(rt)
    print(f"N={n}: {rt:.4f} s")

# Plot N vs runtime
plt.figure(figsize=(7,5))
plt.plot(sizes, runtimes, marker='o')
plt.xlabel("Number of tasks (N)")
plt.ylabel("Runtime (seconds)")
plt.show()
```

```
Input size scaling (avg of 10 runs each task):
N=10: 0.0002 s
N=50: 0.0007 s
N=100: 0.0014 s
N=200: 0.0027 s
N=300: 0.0033 s
N=400: 0.0043 s
N=500: 0.0058 s
N=600: 0.0081 s
N=700: 0.0106 s
N=800: 0.0139 s
N=900: 0.0193 s
N=1000: 0.0209 s
N=1200: 0.0285 s
N=1400: 0.0379 s
N=1600: 0.0479 s
N=1800: 0.0609 s
N=2000: 0.0726 s
N=2200: 0.0879 s
N=2400: 0.1022 s
N=2600: 0.1227 s
```





```

: import numpy as np

logN = np.log2(sizes) # log base-2
n2   = np.array(sizes, dtype=float)**2

scale_log = runtimes[-1] / logN[-1]
scale_n2  = runtimes[-1] / n2[-1]
logN_scaled = logN * scale_log
n2_scaled = n2 * scale_n2

# Re-plot with theory curves
plt.figure(figsize=(7,5))
plt.plot(sizes, runtimes, marker='o', label="Empirical runtime")
plt.plot(sizes, logN_scaled, '--', label="log N (scaled)")
plt.plot(sizes, n2_scaled, ':', label="N^2 (scaled)")
plt.xlabel("Number of tasks (N)")
plt.ylabel("Runtime (seconds)")
plt.grid(True, linestyle="--", alpha=0.6)
plt.legend()
plt.show()

```

